

LLM-Based Generation of BPMN Workflows From Textual Descriptions


NIVON Quentin, SALAÜN Gwen




What is BPMN?



- A **workflow-based notation** created in 2004 by the Business Process Management Initiative (BPMP) and the Object Management Group (OMG).
- It aims at **representing business processes** in a way that is **understandable for both experienced and novice users**.
- An **ISO/IEC standard** since version 2.0 in 2013.



Sequence
Flow




Task




Annotation


Association



Initial
Event



End
Event




Event-based
Gateway



Group

etc.



Exclusive
Gateway



Parallel
Gateway




Inclusive
Gateway

Message flow




Message task




Supported




Sequence
Flow




Task



Initial
Event



End
Event




Exclusive
Gateway



Parallel
Gateway



Annotation



Event-based
Gateway



Inclusive
Gateway

.....
Association




Group

etc.

Message flow




Message task




- Companies are making use of the BPMN notation to **represent their business processes**.
- They **hire experts** to analyse and design the **most adequate BPMN process** according to their needs.
- These processes are often **syntactically/semantically incorrect**.

- What if you do not know **how to write BPMN**?
- What if you do not want to **spend time designing** your BPMN process graphically?
- How can you be sure that your BPMN process is **syntactically/semantically correct**?






General Solution



Internal Computations

General Solution



Internal Computations




First of all, an employee
CollectGoods. Then, the client
PayForDelivery while the
employee PrepareParcel.
Finally, the company can
either DeliverByCar or
DeliverByDrone (depending
on the distance for example)

Textual Representation of the Process

First of all, an employee CollectGoods. Then, the client PayForDelivery while the employee PrepareParcel. Finally, the company can either DeliverByCar or DeliverByDrone (depending on the distance for example)

**Textual Representation
of the Process**



**Large Language
Model (LLM)**

Overview of our Solution

First of all, an employee CollectGoods. Then, the client PayForDelivery while the employee PrepareParcel. Finally, the company can either DeliverByCar or DeliverByDrone (depending on the distance for example)

**Textual Representation
of the Process**



**Large Language
Model (LLM)**

- CollectGoods < (PayForDelivery, PrepareParcel)
- (PayForDelivery, PrepareParcel) < (DeliverByCar, DeliverByDrone)

$\langle E \rangle ::= \mathfrak{t} \mid (\langle E \rangle) \mid$
 $\langle E_1 \rangle \langle \text{op} \rangle \langle E_2 \rangle \mid (\langle E_1 \rangle)^*$
 $\langle \text{op} \rangle ::= ' \mid \& \mid '<' \mid ','$

**Expressions Following
an Internal Grammar**

Overview of our Solution

First of all, an employee CollectGoods. Then, the client PayForDelivery while the employee PrepareParcel. Finally, the company can either DeliverByCar or DeliverByDrone (depending on the distance for example)

**Textual Representation
of the Process**




**Large Language
Model (LLM)**

- CollectGoods < (PayForDelivery, PrepareParcel)
- (PayForDelivery, PrepareParcel) < (DeliverByCar, DeliverByDrone)

$\langle E \rangle ::= \mathbf{t} \mid (\langle E \rangle) \mid$
 $\langle E_1 \rangle \langle \text{op} \rangle \langle E_2 \rangle \mid (\langle E_1 \rangle)^*$
 $\langle \text{op} \rangle ::= \text{'|'} \mid \text{'\&'} \mid \text{'<'} \mid \text{'\;'}$

**Expressions Following
an Internal Grammar**



Abstract Syntax Trees

Overview of our Solution

First of all, an employee CollectGoods. Then, the client PayForDelivery while the employee PrepareParcel. Finally, the company can either DeliverByCar or DeliverByDrone (depending on the distance for example)

**Textual Representation
of the Process**




**Large Language
Model (LLM)**


- CollectGoods < (PayForDelivery, PrepareParcel)
- (PayForDelivery, PrepareParcel) < (DeliverByCar, DeliverByDrone)

$\langle E \rangle ::= \mathbf{t} \mid (\langle E \rangle) \mid$
 $\langle E_1 \rangle \langle \text{op} \rangle \langle E_2 \rangle \mid (\langle E_1 \rangle)^*$
 $\langle \text{op} \rangle ::= \text{'|'} \mid \text{'\&'} \mid \text{'<'} \mid \text{'\,'}$

**Expressions Following
an Internal Grammar**



Dependency Graph



Abstract Syntax Trees

Overview of our Solution

First of all, an employee CollectGoods. Then, the client PayForDelivery while the employee PrepareParcel. Finally, the company can either DeliverByCar or DeliverByDrone (depending on the distance for example)

Textual Representation of the Process




Large Language Model (LLM)


- CollectGoods < (PayForDelivery, PrepareParcel)
- (PayForDelivery, PrepareParcel) < (DeliverByCar, DeliverByDrone)

$\langle E \rangle ::= \mathbf{t} \mid (\langle E \rangle) \mid \langle E_1 \rangle \langle \text{op} \rangle \langle E_2 \rangle \mid (\langle E_1 \rangle)^*$
 $\langle \text{op} \rangle ::= \text{'|'} \mid \text{'&'} \mid \text{'<'} \mid \text{';'}$


Expressions Following an Internal Grammar



BPMN Process



Dependency Graph



Abstract Syntax Trees

The user first has to write a **textual description** of the process-to-be.

First, the banker either CreateProfile (CP) for the user, or, if it is not needed, he RetrieveCustomerProfile (RCP) which triggers the system to perform the AnalyseCustomerProfile (ACP) task.

Then, the user executes the task ReceiveSupportDocuments (RSD) so that the system can start UpdateInfoRecords (UID) and perform a BackgroundVerification (BV).

If the verification finds missing or incorrect information, the system RequestAdditionalInfo (RAI) to the user, who has to ReceiveSupportDocuments (RSD) again.

Otherwise, the process ends with CreateAccount (CA).

The textual description is then **given to a (fine-tuned) LLM** (GPT 3.5 atm).

First, the banker either CreateProfile (CP) for the user, or, if it is not needed, he RetrieveCustomerProfile (RCP) which triggers the system to perform the AnalyseCustomerProfile (ACP) task. Then, the user executes the task ReceiveSupportDocuments (RSD) so that the system can start UpdateInfoRecords (UID) and perform a BackgroundVerification (BV). If the verification finds missing or incorrect information, the system RequestAdditionalInfo (RAI) to the user, who has to ReceiveSupportDocuments (RSD) again. Otherwise, the process ends with CreateAccount (CA).



The textual description is then **given to a (fine-tuned) LLM** (GPT 3.5 atm).

First, the banker either CreateProfile (CP) for the user, or, if it is not needed, he RetrieveCustomerProfile (RCP) which triggers the system to perform the AnalyseCustomerProfile (ACP) task. Then, the user executes the task ReceiveSupportDocuments (RSD) so that the system can start UpdateInfoRecords (UID) and perform a BackgroundVerification (BV). If the verification finds missing or incorrect information, the system RequestAdditionalInfo (RAI) to the user, who has to ReceiveSupportDocuments (RSD) again. Otherwise, the process ends with CreateAccount (CA).



The LLM processes the description and returns a **set of expressions** following an **internal grammar**.

$$\begin{aligned}\langle E \rangle &::= \mathbf{t} \quad | \quad (\langle E \rangle) \quad | \quad \langle E_1 \rangle \langle \text{op} \rangle \langle E_2 \rangle \quad | \quad (\langle E_1 \rangle)^* \\ \langle \text{op} \rangle &::= \text{'|'} \quad | \quad \text{'\&'} \quad | \quad \text{'<'} \quad | \quad \text{';'}\end{aligned}$$

Given our description, the LLM returns **three expressions**:

(RetrieveCustomerProfile < AnalyseCustomerProfile) | CreateProfile

Given our description, the LLM returns **three expressions**:

(RetrieveCustomerProfile < AnalyseCustomerProfile) | CreateProfile

(RetrieveCustomerProfile, AnalyseCustomerProfile, CreateProfile) <
(ReceiveSupportDocuments < (UpdateInfoRecords, BackgroundVerification))

Given our description, the LLM returns **three expressions**:




(RetrieveCustomerProfile < AnalyseCustomerProfile) | CreateProfile

(RetrieveCustomerProfile, AnalyseCustomerProfile, CreateProfile) <
(ReceiveSupportDocuments < (UpdateInfoRecords, BackgroundVerification))


(UpdateInfoRecords, BackgroundVerification) < ((RequestAdditionalInfo <
ReceiveSupportDocuments) | CreateAccount)

Detailed Solution – Step 4

These expressions are then **mapped to** their corresponding **abstract syntax trees (ASTs)**.




The **sequential information** contained in the multiple ASTs is gathered to obtain a **cleaner** representation of it, called **dependency graph**.



Detailed Solution – Step 6


If the dependency graph **contains loops**, they are **analysed**, and all the **information needed to reconstruct** them is extracted.




Entry node(s) computation

Detailed Solution – Step 6

If the dependency graph **contains loops**, they are **analysed**, and all the **information needed to reconstruct** them is extracted.




Entry node(s) computation




Exit node(s) computation

Detailed Solution – Step 6


If the dependency graph **contains loops**, they are **analysed**, and all the **information needed to reconstruct** them is extracted.



Entry node(s) computation




Exit node(s) computation




Mandatory path(s) computation

Detailed Solution – Step 6


If the dependency graph **contains loops**, they are **analysed**, and all the **information needed to reconstruct** them is extracted.



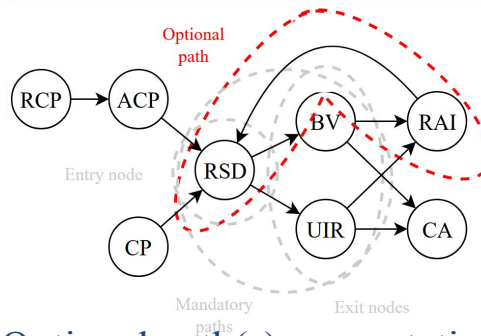
Entry node(s) computation



Exit node(s) computation




Mandatory path(s) computation



Optional path(s) computation


Detailed Solution – Step 7

Once the loops have been retrieved, the **AST corresponding to the dependency graph** is built from the dependency graph.



Detailed Solution – Step 8

The **ultimate step** to obtain an AST containing all the information belonging to the original expressions consists in **inserting the choices**.

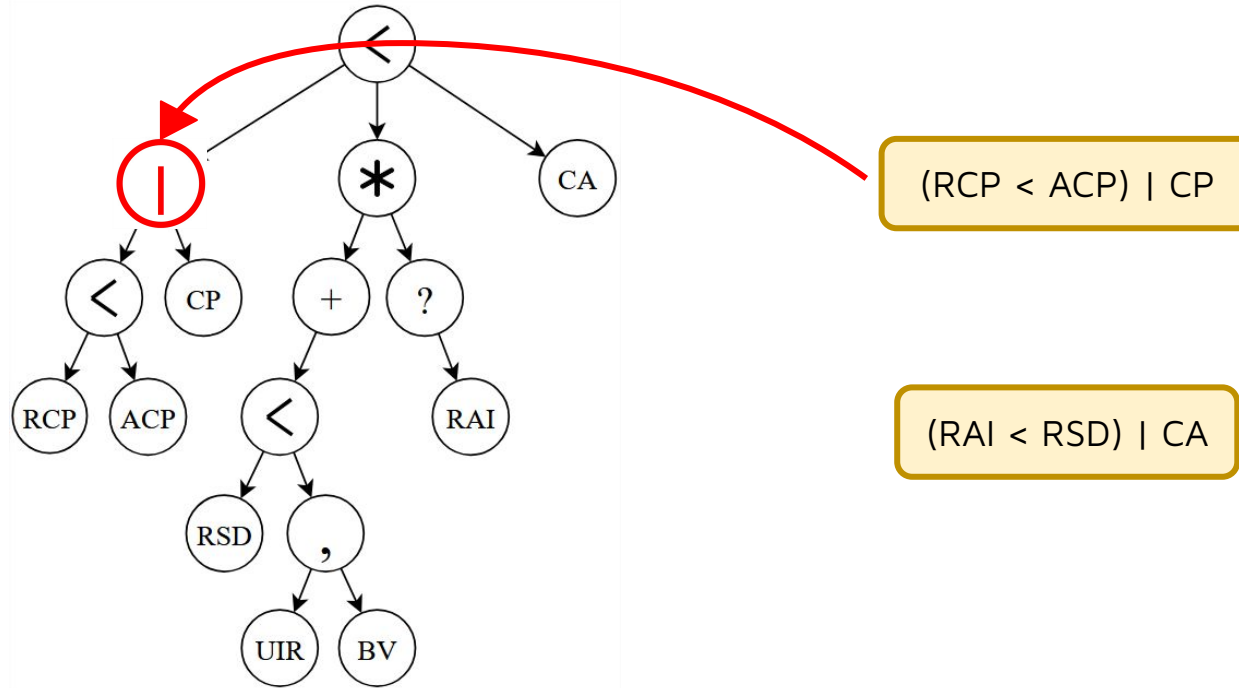


$(RCP < ACP) | CP$

$(RAI < RSD) | CA$

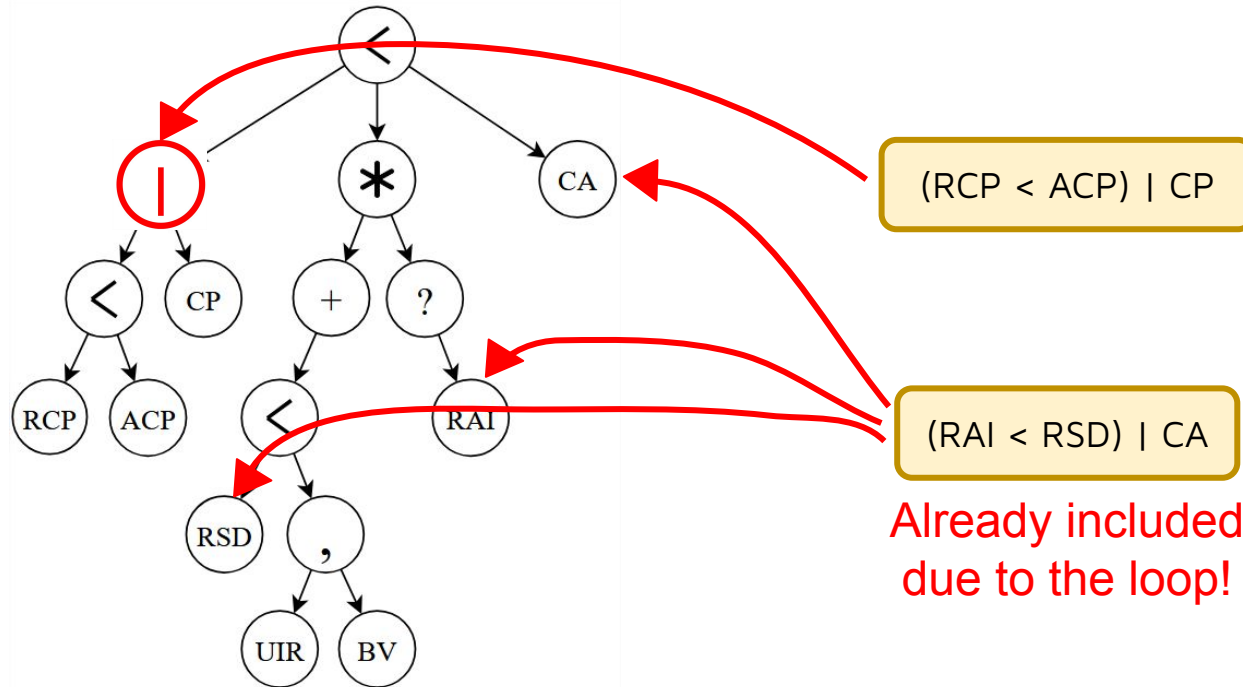
Detailed Solution – Step 8

The **ultimate step** to obtain an AST containing all the information belonging to the original expressions consists in **inserting the choices**.



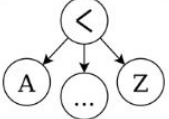

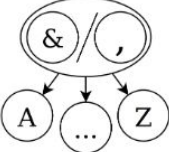
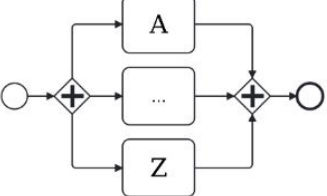
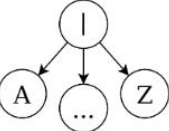
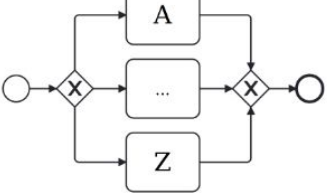
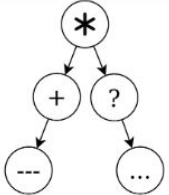
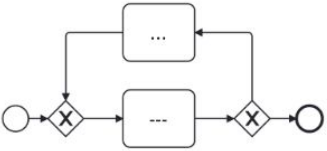
Detailed Solution – Step 8

The **ultimate step** to obtain an AST containing all the information belonging to the original expressions consists in **inserting the choices**.

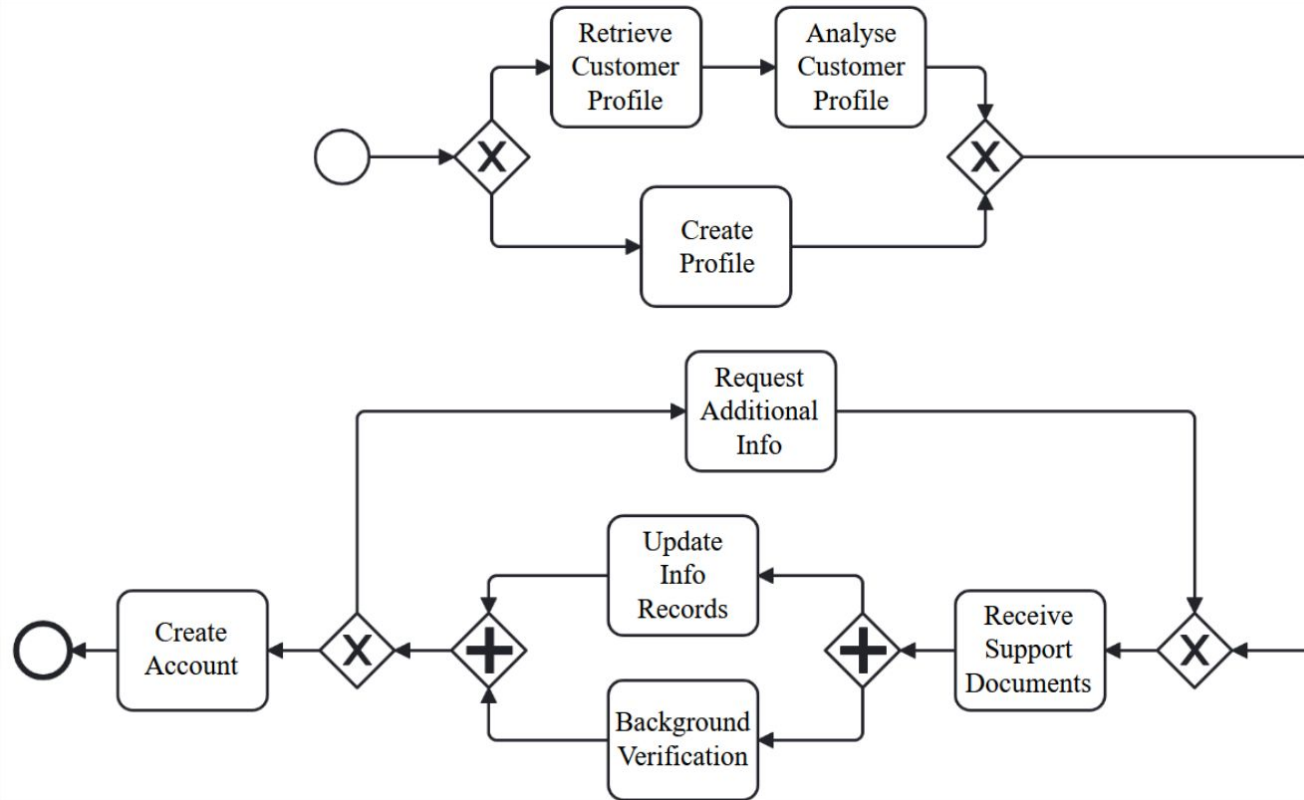


Now that the AST contains the choices, the **BPMN process is ready to be generated**.

To do so, **patterns** are applied recursively to the merged AST, **starting from the deepest nodes** (leafs).

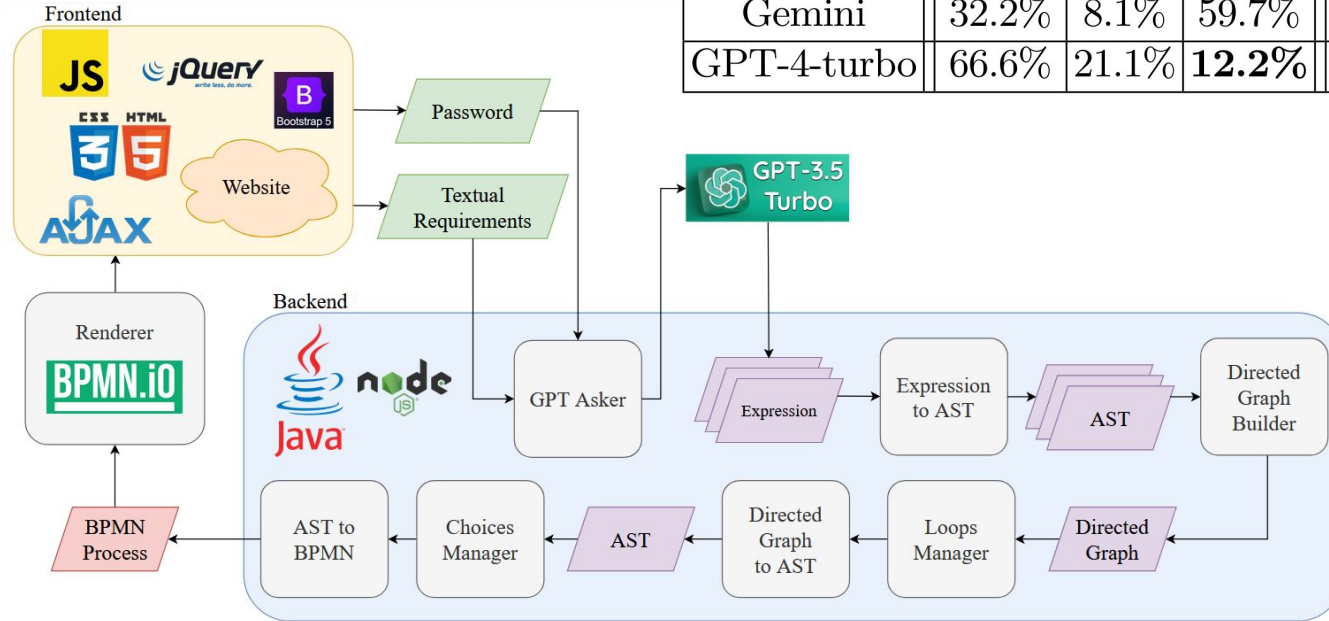
Pattern	AST	BPMN
(1) Sequential Pattern		
(2) Parallel Pattern		
(3) Choice Pattern		
(4) Loop Pattern		


Detailed Solution – Result



Tool – Presentation & Experiments

Tool/Model	✓	?	✗	Avg. Exec. Time (s)
Our tool	78.5%	8%	13.5%	4.07
ProMoAI	50%	8.7%	41.2%	24.7
Gemini	32.2%	8.1%	59.7%	8.32
GPT-4-turbo	66.6%	21.1%	12.2%	19.2





Link: <https://quentinnivon.github.io/pages/givup.html>

In this work, we proposed a **9 steps approach** aiming at automatically designing **syntactically and semantically correct BPMN** processes from a **textual description** of the requirements.

This work offers several **perspectives**:

- Get rid of ASTs to allow more complex constructs (intricate loops/unbalanced gateways) } [TSE'25]
- Formalise the transformation operations
- Add new features (such as model checking of textual properties) } [FSE'25]
- Adapt the description to (visual) process changes } David's work